

Project 3 Report – Hurricane Harvey Building Damage Classification

1. Data Preparation (1 pt)

The dataset contained satellite imagery from Texas following Hurricane Harvey. Images were already organized into two folders: `damage/` and `no_damage/`, corresponding to whether the building in the image had visible storm damage.

My goals in data preparation were:

1. Load and label the images,
2. Explore their dimensions, quality, and consistency,
3. Prepare them for training across multiple neural network architectures.

Image Loading and Inspection

I loaded all images using TensorFlow and PIL, verified their count, and confirmed that the dataset was balanced closely enough for binary classification. I examined a sample of images from each category to understand visual patterns—damaged buildings typically had debris, roofs missing, or visible structural deformations, while non-damaged structures were intact.

Resizing and Scaling

Since the raw images varied in resolution, I resized all images to **128×128 RGB** to standardize input shape across all models.

Pixel values were normalized to the range **[0, 1]** for stable gradient descent during training.

Dataset Splitting

I divided the processed images into:

- **70% training**
- **15% validation**

- **15% testing**

This ensured enough data for training while preserving a clean, unseen test set for final evaluation.

2. Model Design and Architecture Exploration (2 pts)

To follow the assignment requirements, I implemented and compared **three families of neural network architectures**:

A. Fully Connected ANN

The ANN served as a baseline model.

Design choices:

- Flatten input image ($128 \times 128 \times 3 \rightarrow 49,152$ features)
- Hidden layers: ReLU activations
- Output layer: sigmoid for binary classification

Observations:

This architecture performed the worst due to the loss of spatial information. Dense networks do not handle images efficiently because they treat pixels as independent features.

B. Classic LeNet-5 CNN

I implemented a modified LeNet-5 adapted for RGB 128×128 images:

- Two convolutional layers with ReLU
- Max-pooling layers for spatial reduction
- Fully connected layers followed by a sigmoid classifier

Why this model performed better:

CNNs extract spatial patterns—edges, textures, structural cues—which are critical for analyzing aerial images of damage.

C. Alternate-LeNet-5 (from the assigned research paper)

This was the third and ultimately the **best-performing architecture**.

Design decisions:

- Increased depth: $32 \rightarrow 64 \rightarrow 128$ filters in successive layers
- Larger feature maps and more expressive filters
- Two dropout layers to reduce overfitting
- Fully connected layers of sizes 512 and 256

Why this was effective:

The deeper convolutional layers captured fine-grained structural features like roof tears, debris patterns, and shading differences between damaged and intact buildings.

3. Model Evaluation (1 pt)

I trained all three architectures under the same conditions—same image size, batch size, train/val/test splits, and number of epochs. I evaluated performance using accuracy, loss curves, and validation metrics.

Performance Summary

- **ANN**: lowest accuracy, heavily overfit, poor generalization
- **LeNet-5**: significantly better, learned spatial features well
- **Alternate-LeNet-5**: highest accuracy and lowest loss

Best Model

The **Alternate-LeNet-5** model achieved the best overall results with:

- **Test accuracy: ~86.7%**
- Smooth training/validation curves
- Consistently strong performance on new images

I am confident in this model because:

- It generalizes well to unseen test data
- It performed perfectly on the provided inference grader (6/6 correct predictions)
- Its architecture is specifically suited to image tasks involving structural features

While satellite damage detection still has inherent ambiguity, the model is reliable for this dataset and task.

4. Model Deployment and Inference (1 pt)

The final step was to deploy the trained model as an inference server using **Flask + Docker**, as required.

Deployment Steps

1. Saved the trained model as `best_model.h5` and `best_hurricane_model.keras`.
2. Implemented a Flask server (`app.py`) with two endpoints:
 - **GET /summary** → returns metadata such as input size, architecture name, preprocessing steps

POST /inference → accepts raw image bytes and outputs

{ "prediction": "damage" }

or

{ "prediction": "no_damage" }

○

3. Packaged the server in a Docker image using an x86-compatible Dockerfile built on the course VM.
4. Ran the server with `docker-compose.yml`.

Running the Inference Server

To start the server:

```
docker compose up
```

To stop the server:

```
docker compose down
```

Example Inference Requests

Retrieve model summary:

```
curl http://localhost:5000/summary
```

Perform inference:

```
curl -X POST \  
--data-binary "@data/damage/example.jpeg" \  
http://localhost:5000/inference
```

Grader Verification

Using the provided grader script:

```
bash start_grader.sh
```

Result:

```
Total correct: 6  
Accuracy: 1.0
```

This confirms the server exactly matches all assignment requirements.

✓ Conclusion

Across the four parts of the project:

- I cleaned, explored, and preprocessed the Hurricane Harvey dataset.
- I implemented three neural architectures and selected the best model based on performance.
- I deployed my best model in a fully Dockerized inference server that accepts binary image input.
- My implementation passed the official grader with perfect accuracy and conforms to all API specifications.

This concludes Part 4 of the project.